

# Newbie fuzzing gains: Getting the most of your iterations

Presented by @antonio-morales

# Dumb fuzzing

- Random mutations and run cases
- No iterations
- Bruteforce



```
$> while true;\n> do [tested_program] `head -c 100 /dev/urandom;\  
> done
```

*“Coverage-guided”*

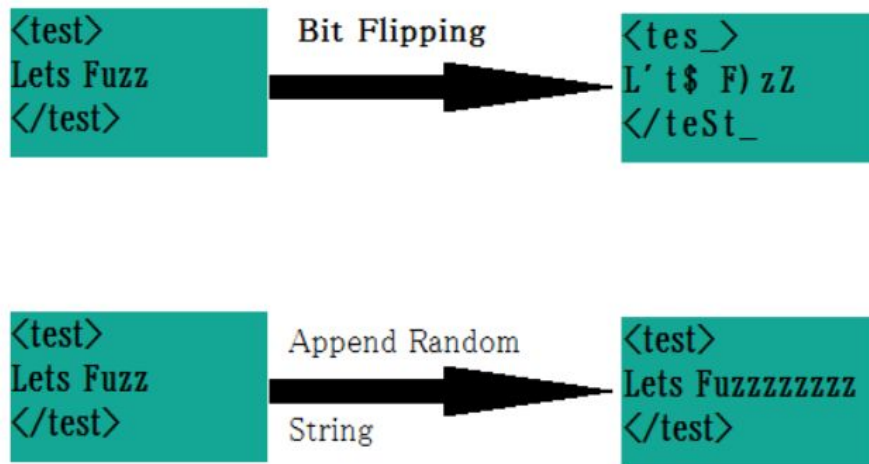
*“Evolutionary”*

*“Mutation-based”*

**FUZZERS**

# Mutation-based fuzzer

- A technique that generates fuzz inputs by applying small mutations to existing valid inputs
- No knowledge of the structure of the input is assumed



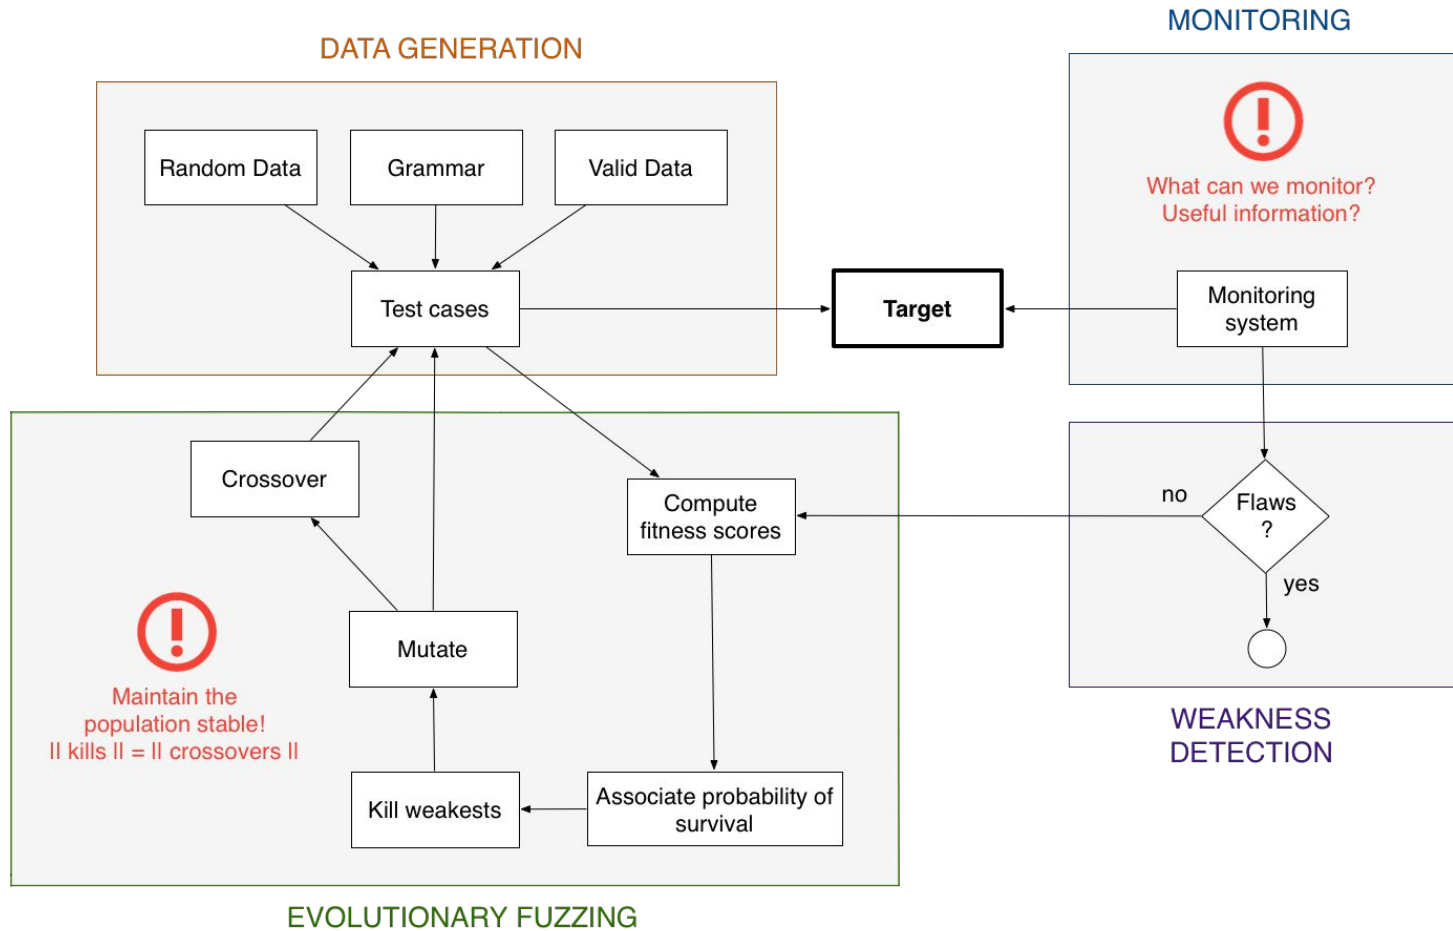
# AFL deterministic mutations

For every byte of the input file:

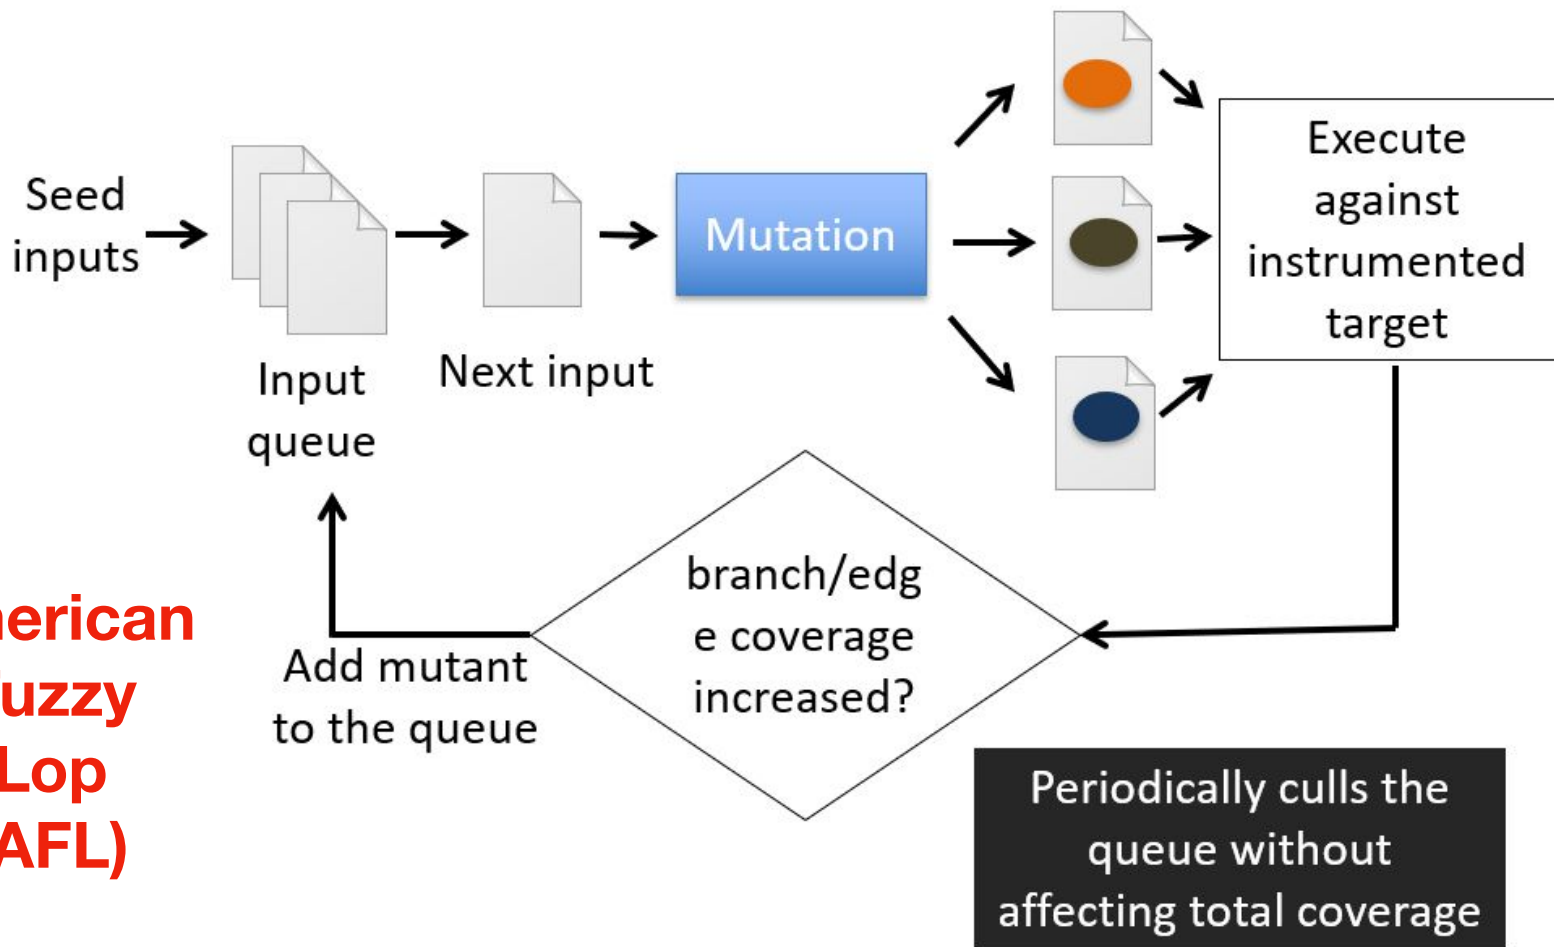
- **Bit flipping:** Sequential and ordered 1,2,4,8,16,32 bit flips.
- **Simple Arithmetics:** increments and decrements of existing integer values (8,16,32 bits integers).
- **Known integers:** Hardcoded set of integer edge cases (-1, 256, 1024, MAX\_INT-1, etc).

# AFL randomized mutations

- **Block deletion**
- **Block duplication** (overwrite or insertion)
- **Block memset**
- **Test case splicing**: involves taking two different input files from the queue that differs in at least two locations and splicing them at a random location in the middle.



# American Fuzzy Lop (AFL)






# Coverage measurement

# Gcov & Lcov:

## Statement/Node coverage

- **Gcov** is a test coverage program that is provided with GCC.
- Gcov analyses the **number of time each line of program** is executed during a run.
- We may then talk of **statement/node coverage**.
- **Lcov** is a graphical front-end for GCC's coverage testing tool gcov.



The screenshot shows a code editor with four tabs: Test.c, coverage\_stubs.c, port.c, and main.c. The Test.c tab is active, displaying C code with line numbers 1 through 29. The code includes a static variable `i` and two functions: `bar` and `TEST_Test`. Statement coverage is indicated by background colors: green for executed statements and red for unexecuted ones. In the `bar` function, lines 12, 13, 15, 17, and 18 are green, while lines 14 and 16 are red. In the `TEST_Test` function, lines 22, 24, 25, 26, and 27 are green, while lines 23 and 28 are red. A vertical scrollbar on the right side of the editor shows the current position within the file.

```
1 7
2 8 #include "Test.h"
3 9
4 10 static int i;
5 11
6 12 static int bar(int i) {
7 13     if (i==0) {
8 14         return 2;
9 15     } else if (i==1) {
10 16         return 0;
11 17     } else {
12 18         return 1;
13 19     }
14 20 }
15 21
16 22 void TEST_Test(int val) {
17 23     if (val==10) {
18 24         i = val;
19 25     } else {
20 26         i += bar(i);
21 27     }
22 28 }
23 29
```

# LCOV - code coverage report

Current view: [top level](#) - /mnt/ex/pseint/pseint

Test: [ads\\_test.info](#)

Date: 2016-03-30 16:15:30

	Hit	Total	Coverage
Lines:	2686	3587	74.9 %
Functions:	199	268	74.3 %

Filename	Line Coverage			Functions	
<a href="#">Ejecutar.cpp</a>		96.9 %	440 / 454	100.0 %	3 / 3
<a href="#">LangSettings.cpp</a>		35.8 %	24 / 67	33.3 %	3 / 9
<a href="#">LangSettings.h</a>		88.6 %	31 / 35	90.0 %	9 / 10
<a href="#">SynCheck.cpp</a>		92.0 %	970 / 1054	90.9 %	20 / 22
<a href="#">case_map.cpp</a>		1.5 %	1 / 67	33.3 %	2 / 6
<a href="#">global.cpp</a>		100.0 %	3 / 3	100.0 %	2 / 2
<a href="#">intercambio.cpp</a>		25.1 %	51 / 203	53.8 %	14 / 26
<a href="#">intercambio.h</a>		100.0 %	5 / 5	85.7 %	6 / 7
<a href="#">main.cpp</a>		46.4 %	115 / 248	80.0 %	4 / 5
<a href="#">new_evaluar.cpp</a>		85.1 %	428 / 503	100.0 %	16 / 16
<a href="#">new_funciones.cpp</a>		68.0 %	123 / 181	61.3 %	19 / 31
<a href="#">new_funciones.h</a>		100.0 %	23 / 23	100.0 %	9 / 9
<a href="#">new_memoria.cpp</a>		100.0 %	15 / 15	100.0 %	3 / 3
<a href="#">new_memoria.h</a>		89.2 %	157 / 176	97.1 %	34 / 35
<a href="#">new_programa.cpp</a>		100.0 %	1 / 1	100.0 %	2 / 2
<a href="#">new_programa.h</a>		92.3 %	72 / 78	91.3 %	21 / 23
<a href="#">utils.cpp</a>		58.8 %	163 / 277	78.3 %	18 / 23

# More than node coverage

There are other coverage criterias that can be used:

- **Edge coverage:**
  - Has every edge in the CFG been executed?
  - LibFuzzer (native LLVM SanitizerCoverage)
  - AFL (stored in 64kb bitmap)
- Complete path coverage
- All possible states


Edge  
Coverage

Complete path  
coverage

All possible  
states

# Edge coverage

Coverage = { (0,1), (0,2), (1,2),  
(2,3), (2,4), (3,6), (4,5), (4,6), (5,4) }



# More than node coverage

There are other coverage criterias that can be used:

- Edge coverage:
- **Complete path coverage:**
  - Require that all possible execution paths are covered
  - It is not feasible if the graph has a loop
- All possible states


Edge  
Coverage

Complete path  
coverage

All possible  
states

# Complete path coverage

Coverage = { (0,1,2,3,6),  
(0,1,2,4,6), (0,1,2,4,5,4,6),  
(0,1,2,4,5,4,5,4,6), ... }



# More than node coverage

There are other coverage criterias that can be used:

- Edge coverage:
- Complete path coverage
- **All possible states:**
  - All combinations of variable values has been tested
  - Each memory state has been reached

Edge  
Coverage

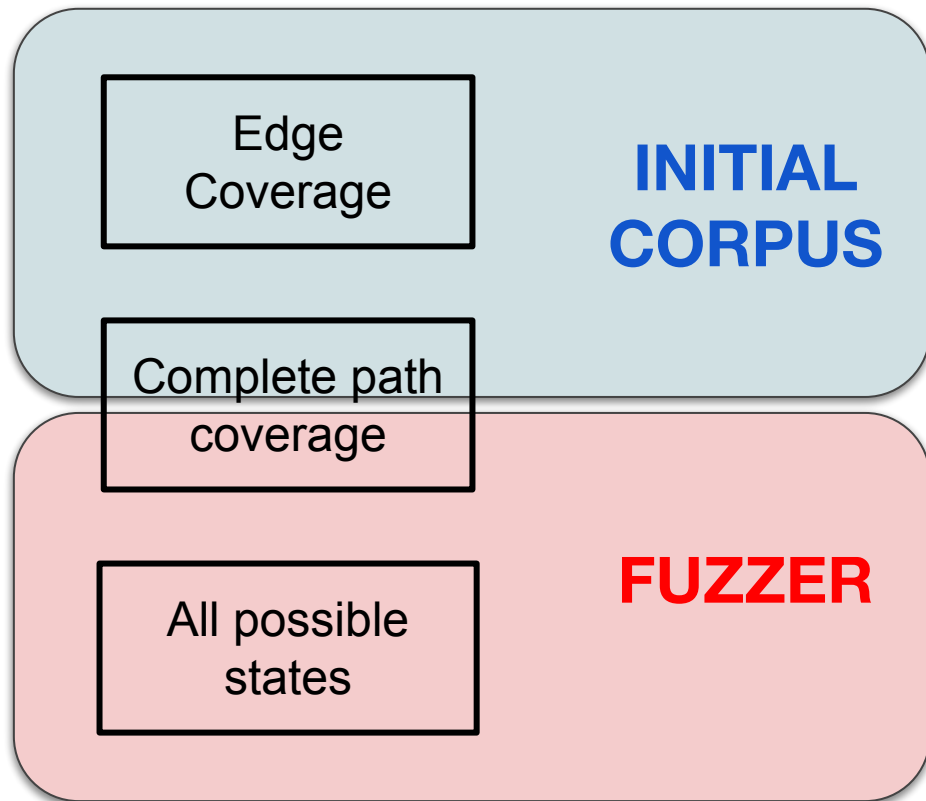
Complete path  
coverage

All possible  
states









# Ideal starting scenario


- Ideally the initial corpus should cover **all code lines** and a **reasonable amount of execution paths** (all functionalities of the program).
- The task of the fuzzer is to find “**Weird paths**” and **unexpected program states**.
- We should also use a **minimal corpus size** and **input size**.



# Size matters

- The **number of mutations** on each stage is correlated with the **size of current input file** (bit/byte mutations)
- Bigger files are usually **more time-consuming**, due to program will have to process more information.

Name	Size ▾	Kind
 HB50 cupcakes.JPG	2 MB	JPEG image
 Roller Skating.JPG	1.3 MB	JPEG image
 50HBJukebox2.jpg	720 KB	JPEG image
 Facebook.tiff	399 KB	TIFF image
 7_days_to_enrol.png	173 KB	PNG image
 JoggingShoes.jpg	71 KB	JPEG image



# Input size minimization

- In the context of a single test case, the following should be maximized:

$$\frac{|program\ states\ explored|}{input\ size}$$

- We should reach the highest byte-to-coverage feature ratio
- AFL: **afl-tmin** tool

```
[+] Read 272 bytes from 'id:002247,src:002216,op:flip2,pos:23.png'.
[*] Performing dry run (mem limit = 0 MB, timeout = 1000 ms)...
[+] Program exits with a signal, minimizing in crash mode.
[*] Stage #0: One-time block normalization...
[+] Block normalization complete, 272 bytes replaced.
[*] --- Pass #1 ---
[*] Stage #1: Removing blocks of data...
    Block length = 32, remaining size = 272
[+] Block removal complete, 272 bytes deleted.
[!] WARNING: Down to zero bytes - check the command line and mem limit!
[*] Stage #2: Minimizing symbols (0 code points)...
[+] Symbol minimization finished, 0 symbols (0 bytes) replaced.
[*] Stage #3: Character minimization...
[+] Character minimization done, 0 bytes replaced.
[*] --- Pass #2 ---
[*] Stage #1: Removing blocks of data...
    Block length = 1, remaining size = 0
[+] Block removal complete, 0 bytes deleted.

    File size reduced by : 100.00% (to 0 bytes)
    Characters simplified : 27200.00%
    Number of execs done : 78
    Fruitless execs : path=0 crash=0 hang=0
```

# Corpus minimization

- Minimized initial corpus is other of the key factors during fuzzing process.
- It should be verified that edge coverage is not reduced (and execution paths!!)
- Libfuzzer: “**merge=1**” option
- AFL: **afl-cmin** tool

```
$ afl-cmin -i Trimmed/ -o Pcap-corpus/ -- tcpdump -ee -vv -nnr @@  
corpus minimization tool for afl-fuzz by <lcantuf@google.com>  
  
[*] Testing the target binary...  
[+] OK, 395 tuples recorded.  
[*] Obtaining traces for input files in 'Trimmed/'...  
    Processing file 1514745/1514745...  
[*] Sorting trace sets (this may take a while)...  
[+] Found 9566 unique tuples across 1514745 files.  
[*] Finding best candidates for each tuple...  
    Processing file 1514745/1514745...  
[*] Sorting candidate list (be patient)...  
[*] Processing candidates and writing output files...  
    Processing tuple 9566/9566...  
[+] Narrowed down to 273 files, saved in 'Pcap-corpus/'.
```

# Overcoming obstacles

# “Stuck” fuzzer

- In this example, we can see that the buggy code is only executed if the variable holds the value “0xabad1dea”.
- This in turn is very unlikely because the value in input is mainly generated by the fuzzer performing byte mutations.
- At this point, the fuzzer needs some help to surpass the conditional statement.

```
if (input == 0xabad1dea) {  
    /* terribly buggy code */  
} else {  
    /* secure code */  
}
```

# Splitting up comparisons

- To overcome this issue we can **split up comparisons into smaller ones** which should guide the fuzzer towards the correct value.
- Now, as soon as the fuzzer guesses the first byte correctly it will execute the nested-if condition and therefore **discover a new path**.
- This will repeat for the other if-statements.

<https://lafintel.wordpress.com/2016/08/15/circumventing-fuzzing-roadblocks-with-compiler-transformations/>

```
if (input == 0xabad1dea) {  
    /* terribly buggy code */  
} else {  
    /* secure code */  
}
```



```
if (input >> 24 == 0xab){  
    if ((input & 0xff0000) >> 16 == 0xad) {  
        if ((input & 0xff00) >> 8 == 0x1d) {  
            if ((input & 0xff) == 0xea) {  
                /* terrible code */  
                goto end;  
            }  
        }  
    }  
}  
  
/* good code */
```



Semmle is joining GitHub



Semmle™

Products ▾


Resources ▾

Company ▾

Blog

Contact

Get started



# VLC vulnerabilities discovered by the Semmle security research team

By: **Semmle Team**

August 19, 2019

Category  
Security

Technical Difficulty  
easy

Reading time  
6 min



Subscribe

Today, the VideoLAN team [announced a new release of VLC](#), fixing 11 vulnerabilities reported by [Antonio Morales Maldonado](#) from the Semmle security research team.

MITRE has issued the following CVE IDs for the vulnerabilities: CVE-2019-14437, CVE-2019-14438, CVE-2019-14438, CVE-2019-14498, CVE-2019-14535, CVE-2019-14534, CVE-2019-14533, CVE-2019-14776, CVE-2019-14778, CVE-2019-14779, CVE-2019-14777, CVE-2019-14970.

## About VLC

The [VLC Media Player](#) (commonly known as just **VLC**) is a popular media player developed by the VideoLAN project. VLC is available on most platforms (Windows, MacOS, Linux, Android, iOS).





```
65  /* Top-Level object */
66  static const vlc_guid_t asf_object_header_guid =
67  {0x75B22630, 0x668E, 0x11CF, {0xA6, 0xD9, 0x00, 0xAA, 0x00, 0x62, 0xCE, 0x6C}};
68
69  static const vlc_guid_t asf_object_data_guid =
70  {0x75B22636, 0x668E, 0x11CF, {0xA6, 0xD9, 0x00, 0xAA, 0x00, 0x62, 0xCE, 0x6C}};
71
72  static const vlc_guid_t asf_object_simple_index_guid =
73  {0x33000890, 0xE5B1, 0x11CF, {0x89, 0xF4, 0x00, 0xA0, 0xC9, 0x03, 0x49, 0xCB}};
74
75  static const vlc_guid_t asf_object_index_guid =
76  {0xD6E229D3, 0x35DA, 0x11D1, {0x90, 0x34, 0x00, 0xA0, 0xC9, 0x03, 0x49, 0xBE}};
77
78  /* Header object */
79  static const vlc_guid_t asf_object_file_properties_guid =
80  {0x8cabdca1, 0xa947, 0x11cf, {0x8e, 0xe4, 0x00, 0xC0, 0x0C, 0x20, 0x53, 0x65}};
81
82  static const vlc_guid_t asf_object_stream_properties_guid =
83  {0xB7DC0791, 0xA9B7, 0x11CF, {0x8E, 0xE6, 0x00, 0xC0, 0x0C, 0x20, 0x53, 0x65}};
84
85  static const vlc_guid_t asf_object_header_extension_guid =
86  {0x5FBF03B5, 0xA92E, 0x11CF, {0x8E, 0xE3, 0x00, 0xC0, 0x0C, 0x20, 0x53, 0x65}};
87
88  static const vlc_guid_t asf_object_codec_list_guid =
89  {0x86D15240, 0x311D, 0x11D0, {0xA3, 0xA4, 0x00, 0xA0, 0xC9, 0x03, 0x48, 0xF6}};
90
```

# Providing a custom dictionary

- Custom dictionaries can be added in order to provide the fuzzer with a list of complex syntax tokens.
- Even when no explicit dictionary is given, afl-fuzz will try to extract existing syntax tokens in the input corpus by watching the instrumentation very closely during deterministic byte flips

```
"\x56\x4C\x43\x52\x4F\x4F\x54\x00"  
"\x30\x26\xB2\x75\x8E\x66\xCF\x11\xA6\xD9\x00\xAA\  
"\x36\x26\xB2\x75\x8E\x66\xCF\x11\xA6\xD9\x00\xAA\  
"\x90\x08\x00\x33\xB1\xE5\xCF\x11\x89\xF4\x00\xA0\  
"\xD3\x29\xE2\xD6\xDA\x35\xD1\x11\x90\x34\x00\xA0\  
"\xA1\xDC\xAB\x8C\x47\xA9\xCF\x11\x8E\xE4\x00\xC0\  
"\x91\x07\xDC\xB7\xB7\xA9\xCF\x11\x8E\xE6\x00\xC0\  
"\xB5\x03\xBF\x5F\x2E\xA9\xCF\x11\x8E\xE3\x00\xC0\  
"\x40\x52\xD1\x86\x1D\x31\xD0\x11\xA3\xA4\x00\xA0\  
"\x01\xCD\x87\xF4\x51\xA9\xCF\x11\x8E\xE6\x00\xC0\  
"\x33\x26\xB2\x75\x8E\x66\xCF\x11\xA6\xD9\x00\xAA\  
"\x40\xA4\xD0\xD2\x07\xE3\xD2\x11\x97\xF0\x00\xA0\  
"\x74\xD4\x06\x18\xDF\xCA\x09\x45\xA4\xBA\x9A\xAB\  
"\xCF\x49\x86\xA0\x75\x47\x70\x46\x8A\x16\x6E\x35\  
"\x5B\xD1\xFE\xD4\xD3\x88\x4F\x45\x81\xF0\xED\x5C\  
"\xEA\xCB\xF8\xC5\xAF\x5B\x77\x48\x84\x67\xAA\x8C\  
"\x40\x9E\x69\xF8\x4D\x5B\xCF\x11\xA8\xFD\x00\x80\  
"\xC0\xEF\x19\xBC\x4D\x5B\xCF\x11\xA8\xFD\x00\x80\  
"\xC0\xCF\xD3\x5B\xE6\x5B\xD0\x11\x73\x7C\x00\x70
```

# CVE-2019-14533

- **Advanced Systems Format (ASF)** is Microsoft's proprietary audio/video container, best known for their most common media types: **WMV** and **WMA**.
- Through the use of a custom dictionary including ASF Object GUIDs, an **Use-After-Free vulnerability** was found in VLC media player. As a result, trying to parse an invalid WMV/ASF (Windows Media Video) file will result in UAF when the video is forwarded.
- This bug could allow an attacker to alter the expected application flow.

# CVE-2019-14533

```
static void DemuxEnd( demux_t *p_demux )
{
    demux_sys_t *p_sys = p_demux->p_sys;

    if( p_sys->p_root )
    {
        ASF_FreeObjectRoot( p_demux->s, p_sys->p_root );
        p_sys->p_root = NULL;
        //p_sys->p_fp should also be nulled
    }

    [...]
}
```

# Dealing with checksums

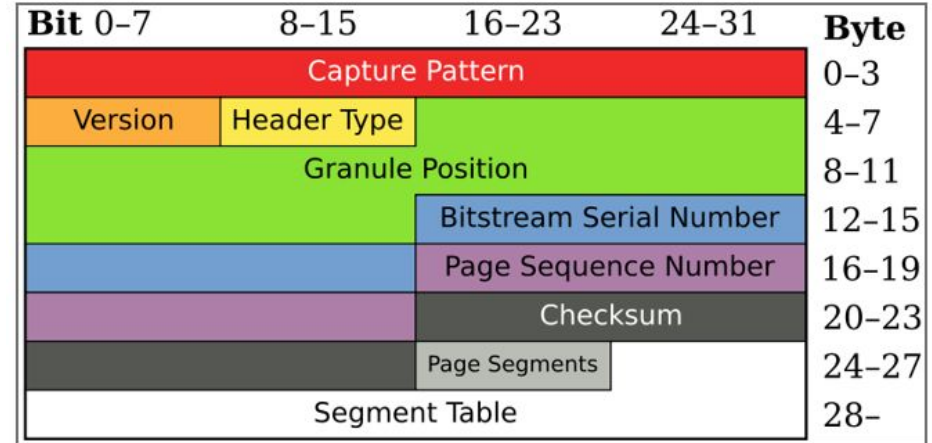
- Some protocols or file formats incorporate checksums that will fail if they're modified arbitrarily.
- There are 2 strategies to counter this:
  - **Re-calculate** the checksum on the fuzzed inputs
  - **Patch** the software to disable checksum tests and repair the checksum fields in malformed inputs

PNG chunk specification.

Name	Size	Description
Length	4 bytes	Length of data field
Type	4 bytes	Chunk type code
Data	n bytes	Data bytes
CRC	4 bytes	CRC of type and data

# CVE-2019-14438

- **Ogg** is a free, open container format that can multiplex a number of independent streams for audio, video, subtitles and metadata.
- OGG page header includes a **CRC32 checksum**, that in the case of VLC, is calculated through **Libogg library**.



# CVE-2019-14438

- After disabling CRC32 checksum we were able to find an **OOB write** in “xiph\_PackHeaders” function.
- As a result is possible to craft the heap through a OOB read/write attack.

```
/* Compare */  
if(memcmp(chksum,page+22,4)){  
    /* D'oh. Mismatch! Corrupt page (or miscapture and not a page  
       at all) */  
    /* replace the computed checksum with the one actually read in */  
    memcpy(page+22,chksum,4);
```

```
for (unsigned i = 0; i < packet_count; i++) {  
    if (packet_size[i] > 0) {  
        memcpy(current, packet[i], packet_size[i]);  
        current += packet_size[i];  
    }  
}
```

**Is just the beginning**



# Is just the beginning

CVE	Type
CVE-2019-14437	OOB read
CVE-2019-14438	OOB write
CVE-2019-14498	Divide-by-zero
CVE-2019-14535	Divide-by-zero
CVE-2019-14534	NULL Pointer Dereference
CVE-2019-14533	Use-After-Free
CVE-2019-14776	OOB read
CVE-2019-14778	Use-After-Free
CVE-2019-14779	OOB Read
CVE-2019-14777	Use-After-Free
CVE-2019-14970	OOB Write

# To be continued...

- In-process fuzzing (persistent mode)
- Custom mutators  
(Structure-Aware Fuzzing)
- Power Schedules
- Mutation Scheduling

**Q & A**